

Anthony Raju Kondaveeti

+91 7993243834 | tonykondaveetijmj98@gmail.com | [LinkedIn](#) | [GitHub](#) | [Website](#)

PROFESSIONAL SUMMARY

AI/ML Engineer with hands-on experience building production-grade RAG systems, multilingual NLP applications, and generative AI solutions. Specialized in designing and building scalable ML pipelines, optimizing LLM-based systems, and deploying end-to-end AI solutions in production environments.

EDUCATION

Christ University <i>Master of Science in Data Science</i>	Bengaluru, India <i>June 2024 – March 2026</i>
Andhra Loyola College <i>Bachelor of Science in Mathematics, Physics, Computer Science; CGPA: 9.02/10</i>	Vijayawada, India <i>June 2020 – May 2023</i>

PROFESSIONAL EXPERIENCE

AI/ML Engineer Intern <i>Demetrix Infotech Pvt. Ltd.</i>	July 2025 – Oct 2025 <i>Bengaluru, India</i>
- Built hybrid RAG system (dense embeddings + BM25 + cross-encoder) achieving 92% contextual accuracy - Improved retrieval precision by 75% through transformer-based re-ranking and prompt optimization - Integrated Mistral AI and Qdrant vector DB for multilingual enterprise document search - Reduced inference latency by 40% via optimized vector indexing and caching	

PUBLICATIONS

I-EAS: Culture-Aware Evaluation of LLM Emotional Intelligence	IEEE ISCS 2025
- Designed computational framework evaluating LLM emotional intelligence across multilingual Indian cultural contexts - DOI: 10.1109/ISCS69371.2025.11385960	

KEY PROJECTS

Narrative Intelligence System – AI Story Analysis Platform <i>Python, spaCy, Mistral AI, Gradio</i>	2026
- Built hybrid NLP system (heuristics + LLM) for analyzing 100–5000 word UGC stories, achieving 92% accuracy in genre classification and engagement scoring - Reduced LLM inference cost by 66% through prompt optimization, context compression, and caching (210ms cached latency) - Implemented evaluation pipeline measuring accuracy, latency, and robustness across multilingual narratives	
CoChain.ai - ML-Powered Student Collaboration Platform <i>Flask, Supabase, RL</i>	2025
- Built full-stack platform enabling 500+ students to discover GitHub projects and form collaboration teams with real-time matching - Implemented Thompson Sampling contextual bandit with SentenceTransformer embeddings improving recommendation relevance by 25% over baseline - Deployed production system with A/B testing framework, rate limiting, and performance monitoring	
YouTube Video Analyzer Chrome Extension <i>Flask, Transformers, JavaScript</i>	2024
- Built Chrome extension with multilingual transcript extraction (English, Hindi, Telugu) and RoBERTa-based sentiment analysis on 300+ comments - Implemented DistilRoBERTa emotion detection with asynchronous processing, YAKE keyword extraction, and timestamp-aligned visualization dashboard	

TECHNICAL SKILLS

Programming Languages: Python, SQL
ML/DL Frameworks: PyTorch, TensorFlow, Keras, Scikit-learn, Hugging Face Transformers
LLM & NLP: LangChain, Sentence-Transformers, Tokenization, Attention Mechanisms, Prompt Engineering, RAG, Cross-Encoder Models, Mistral AI
Vector Databases: Qdrant, ChromaDB
Cloud & Deployment: AWS EC2, Docker, FastAPI, Flask
Evaluation & Experimentation: Precision@k, Recall, BLEU, ROUGE, A/B Testing

CERTIFICATIONS & TEACHING

Certifications: Advanced Learning Algorithms (Coursera) Supervised Machine Learning (Coursera) AWS Cloud Fundamentals SQL for Data Science
Teaching: Python & ML Instructor for 50+ Bio Statistics undergraduates (Christ University, Jan-May 2025)